COLUMBIA UNIVERSITY
IN THE CITY OF NEW YORK

EECS E6893 Big Data Analytics

# Homework #0 Report

Name: Yi Zhang

Uni: yz4140

September 2022

# Question 1

(1) As shown below, the pi is successfully estimated via the cluster called pi and the value is correct.

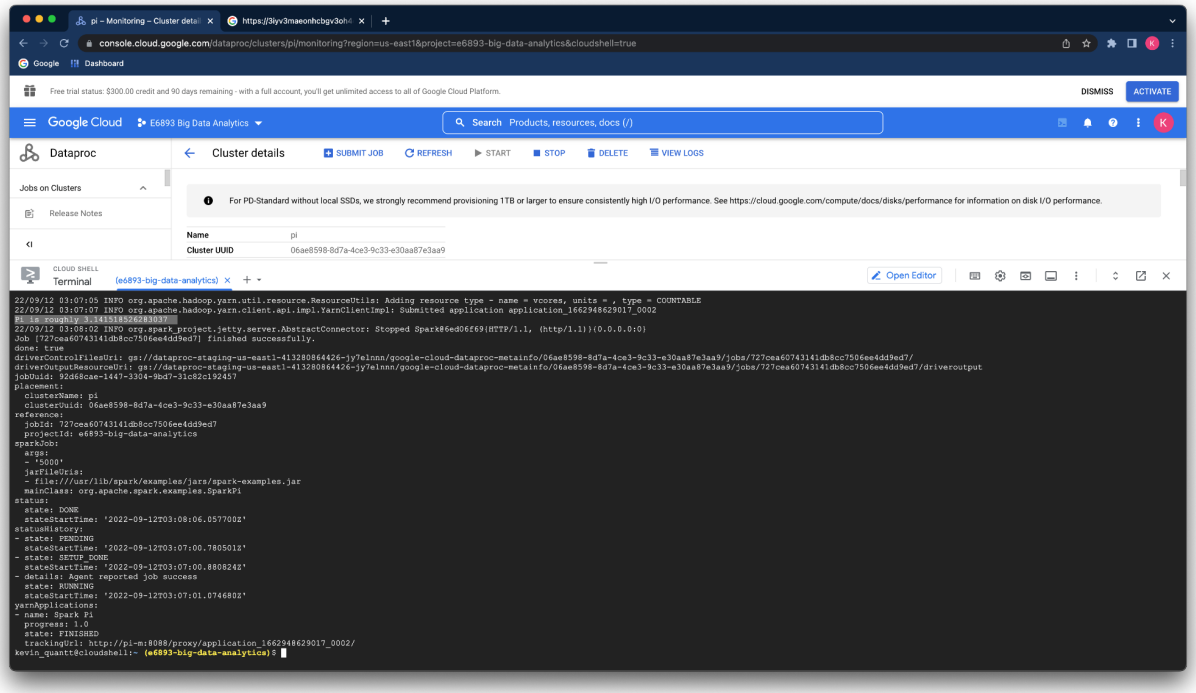


(2) According to the source code provided by the Spark below

Python | Scala | Java

```
def inside(p):
    x, y = random.random(), random.random()
    return x*x + y*y < 1

count = sc.parallelize(range(0, NUM_SAMPLES)) \
             .filter(inside).count()
print("Pi is roughly %f" % (4.0 * count / NUM_SAMPLES))
```

---

The transformation is '**filter**', which helps to create a new dataset containing the (x,y) within the circle. The action is **'count'**, which counts the number of data points in the filtered dataset.

# Question 2

(1) There are 698 stations whose longitude is between -74.04 and -73.94

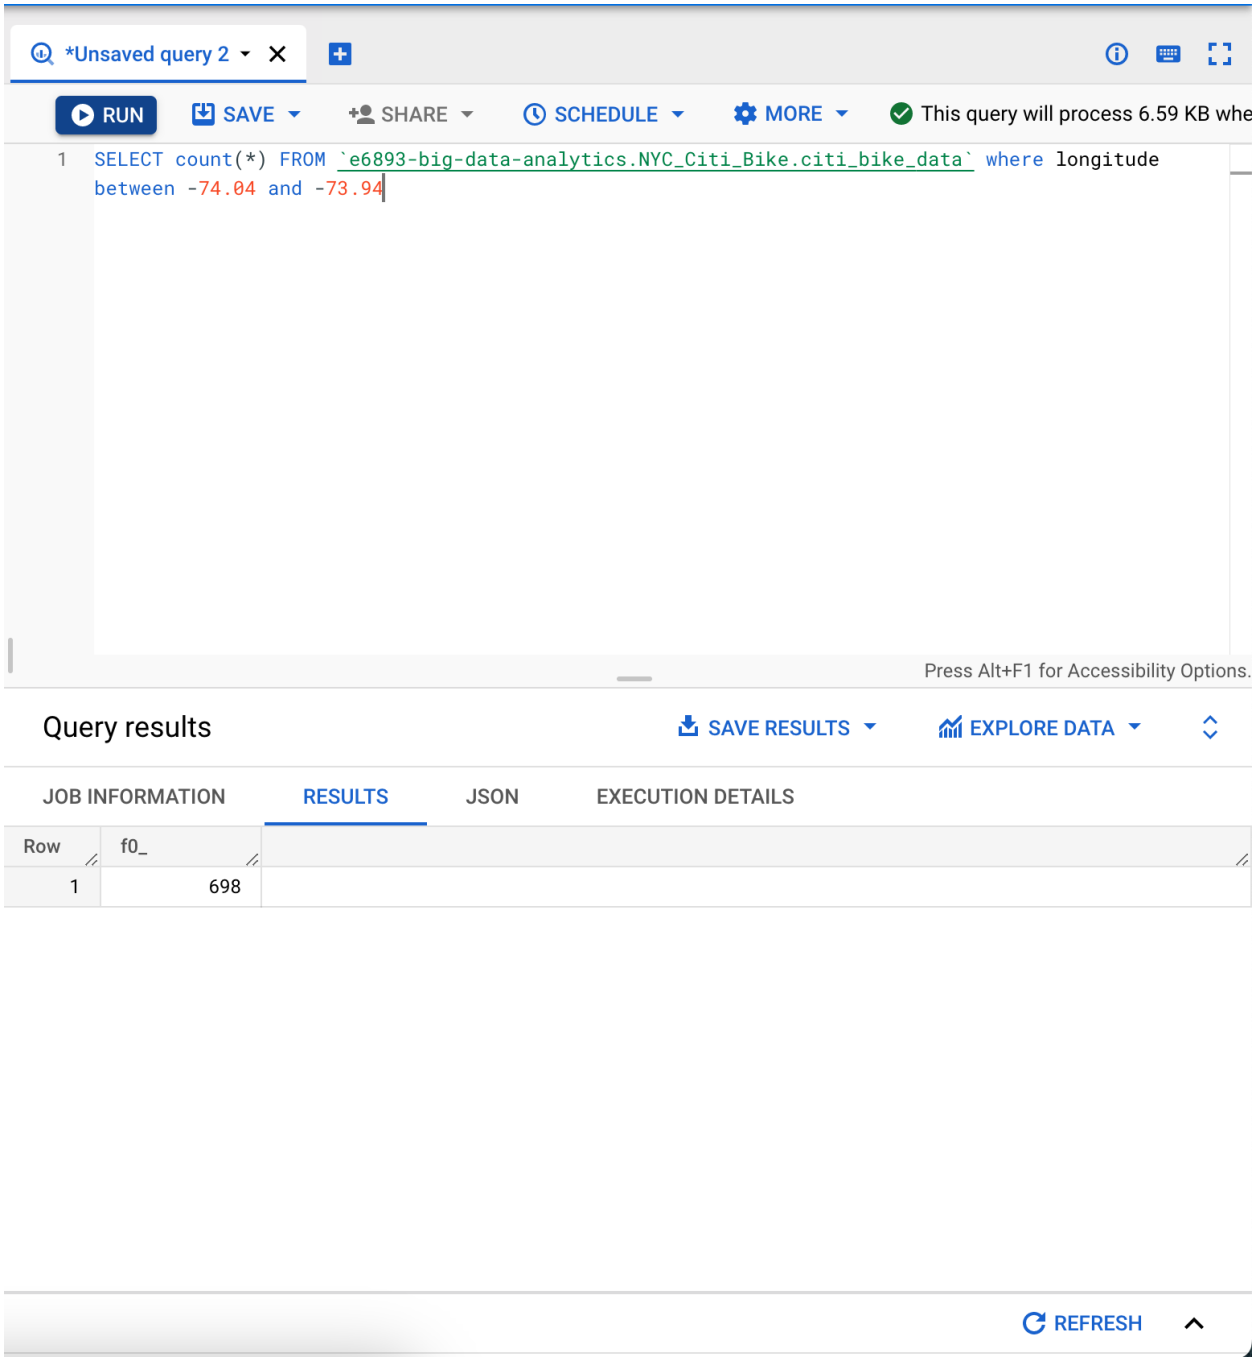


(2) There are 11885 bikes available in region 71

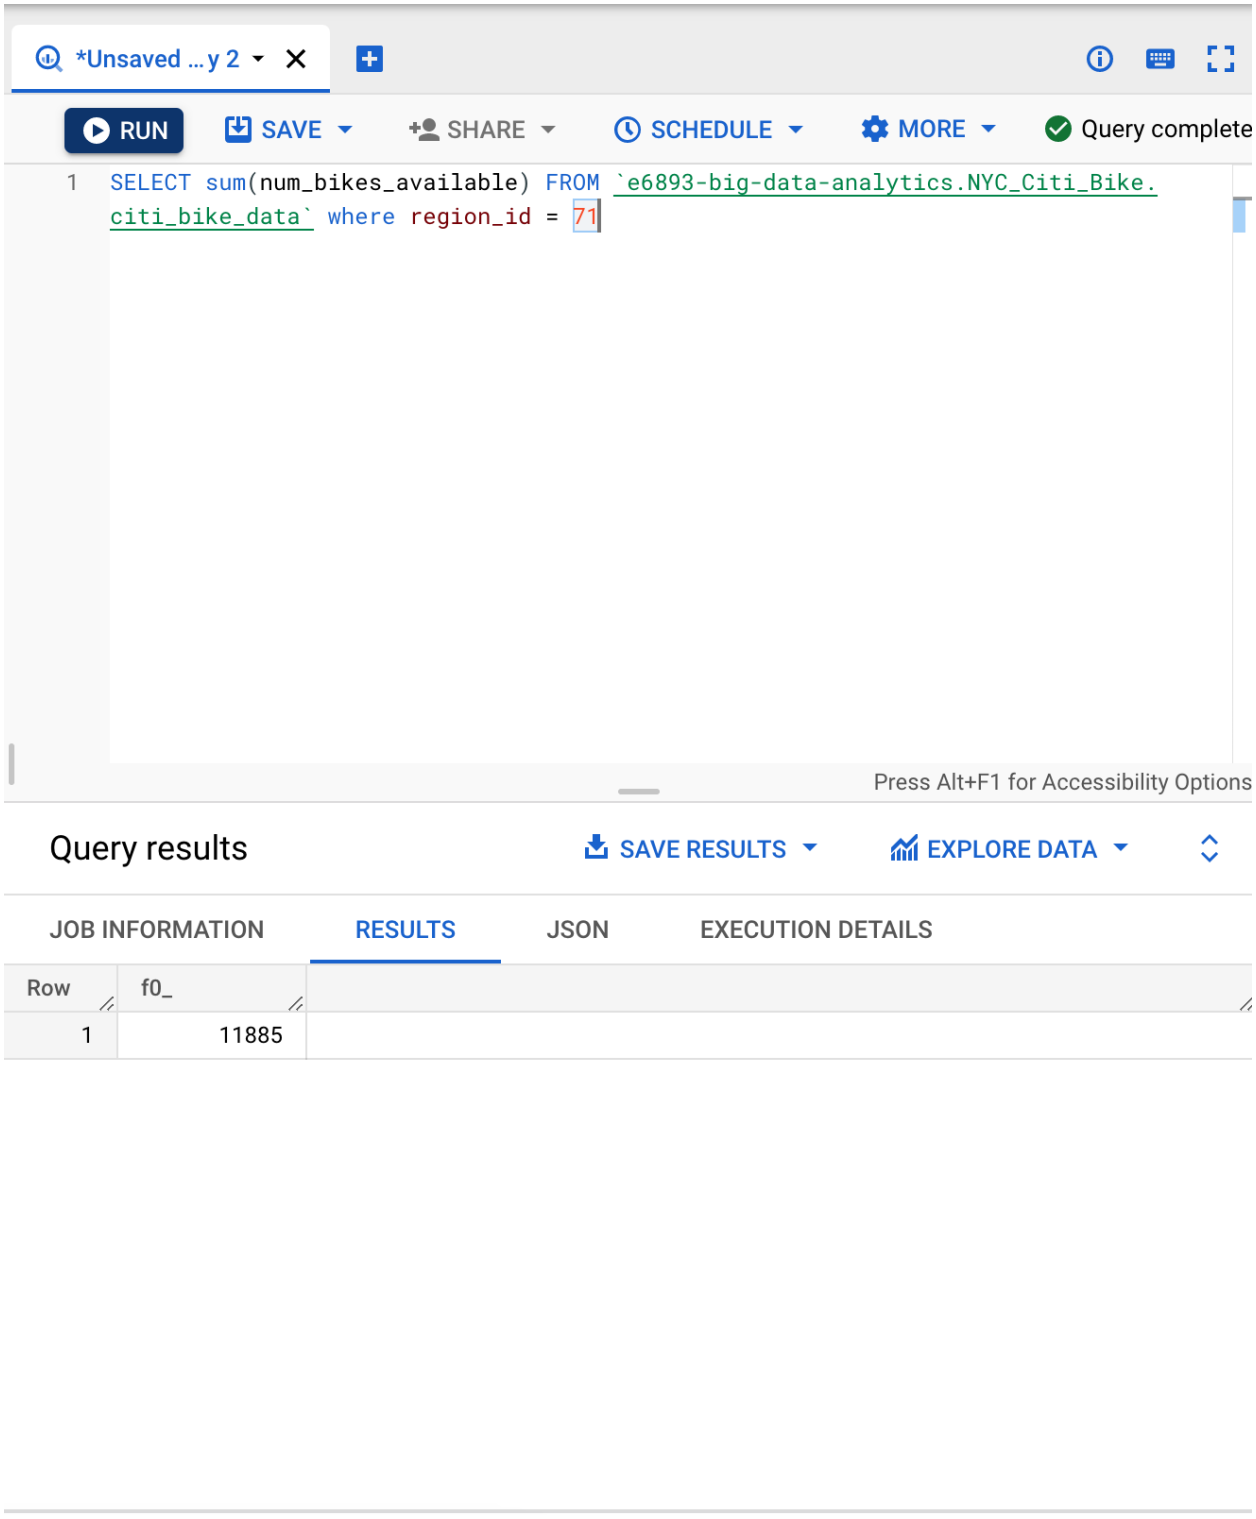


(3) The largest capacity is 79 and station 445, 422, and 501 attain the largest capacity.

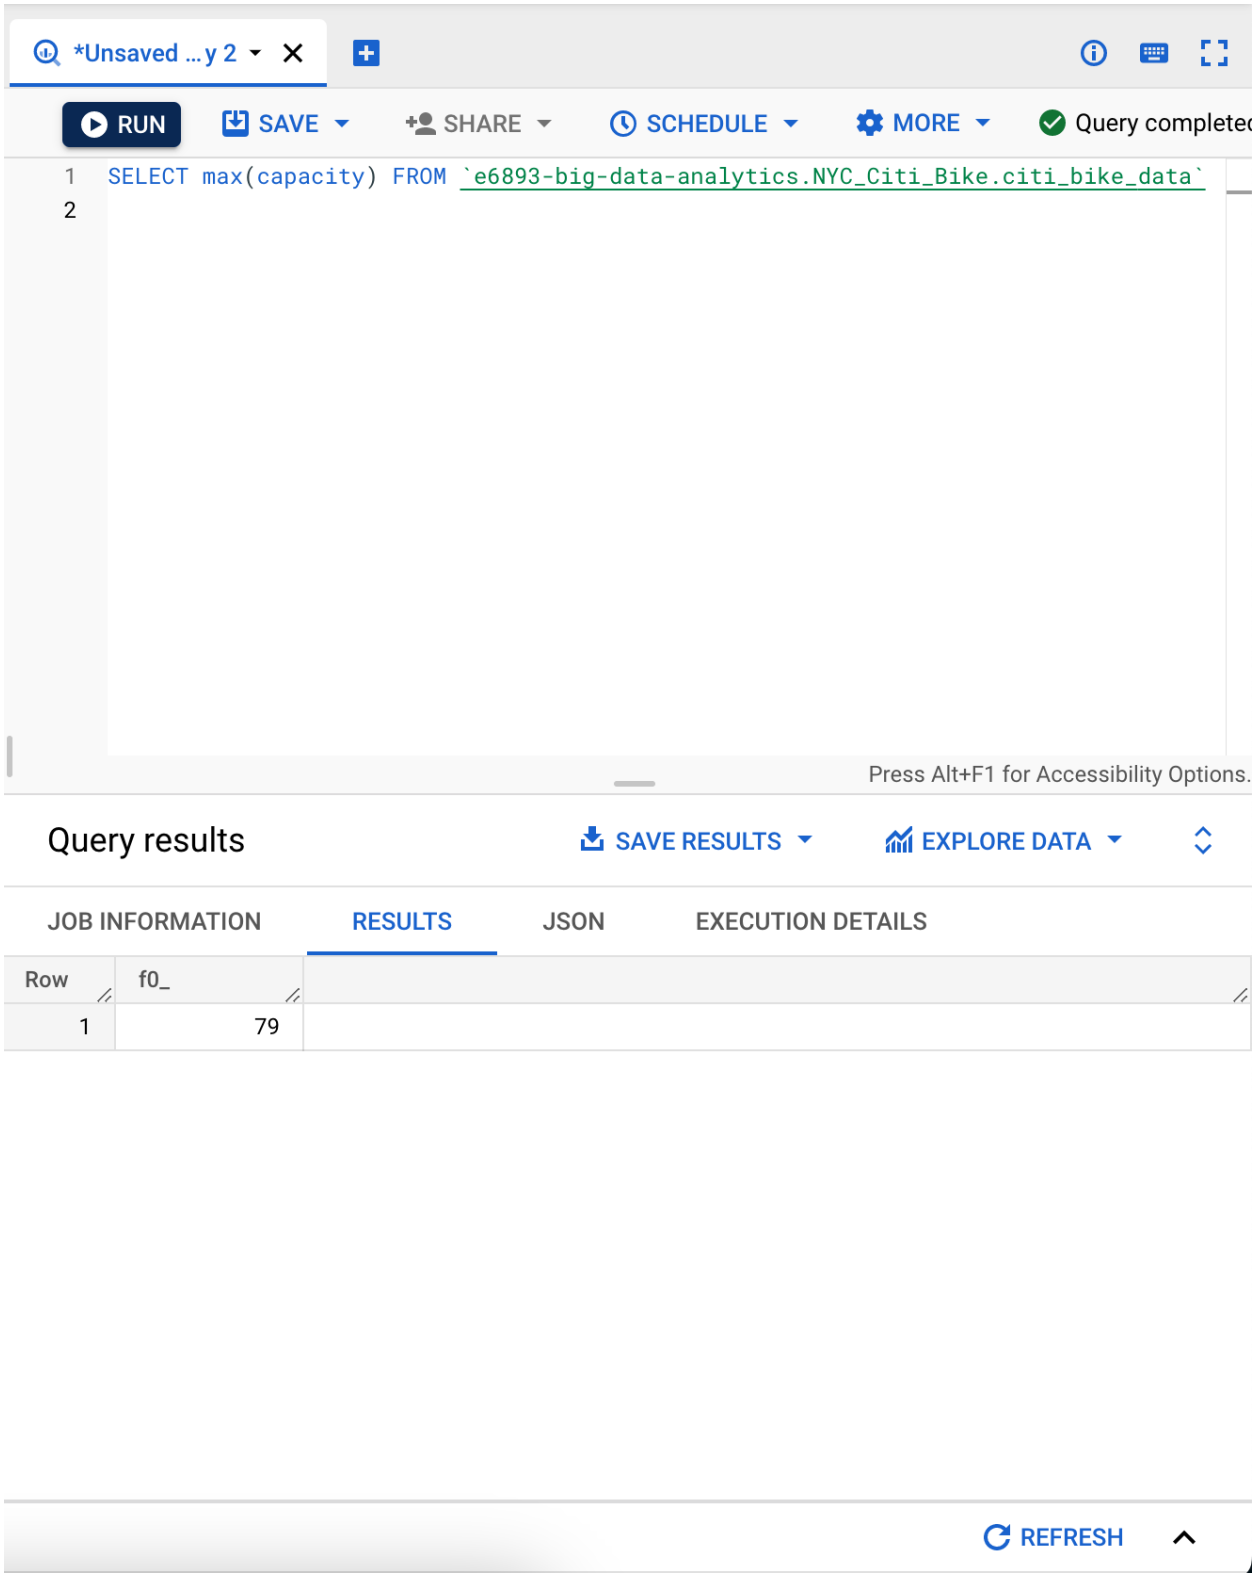
*Unsaved ... y 2
RUN
SAVE
SHARE
SCHEDULE
MORE
Query complete
1 SELECT max(capacity) FROM `e6893-big-data-analytics.NYC_Citi_Bike.citi_bike_data`
2
Press Alt+F1 for Accessibility Options.
Query results
SAVE RESULTS
EXPLORE DATA
JOB INFORMATION
RESULTS
JSON
EXECUTION DETAILS
Row
f0_
1
79
REFRESH

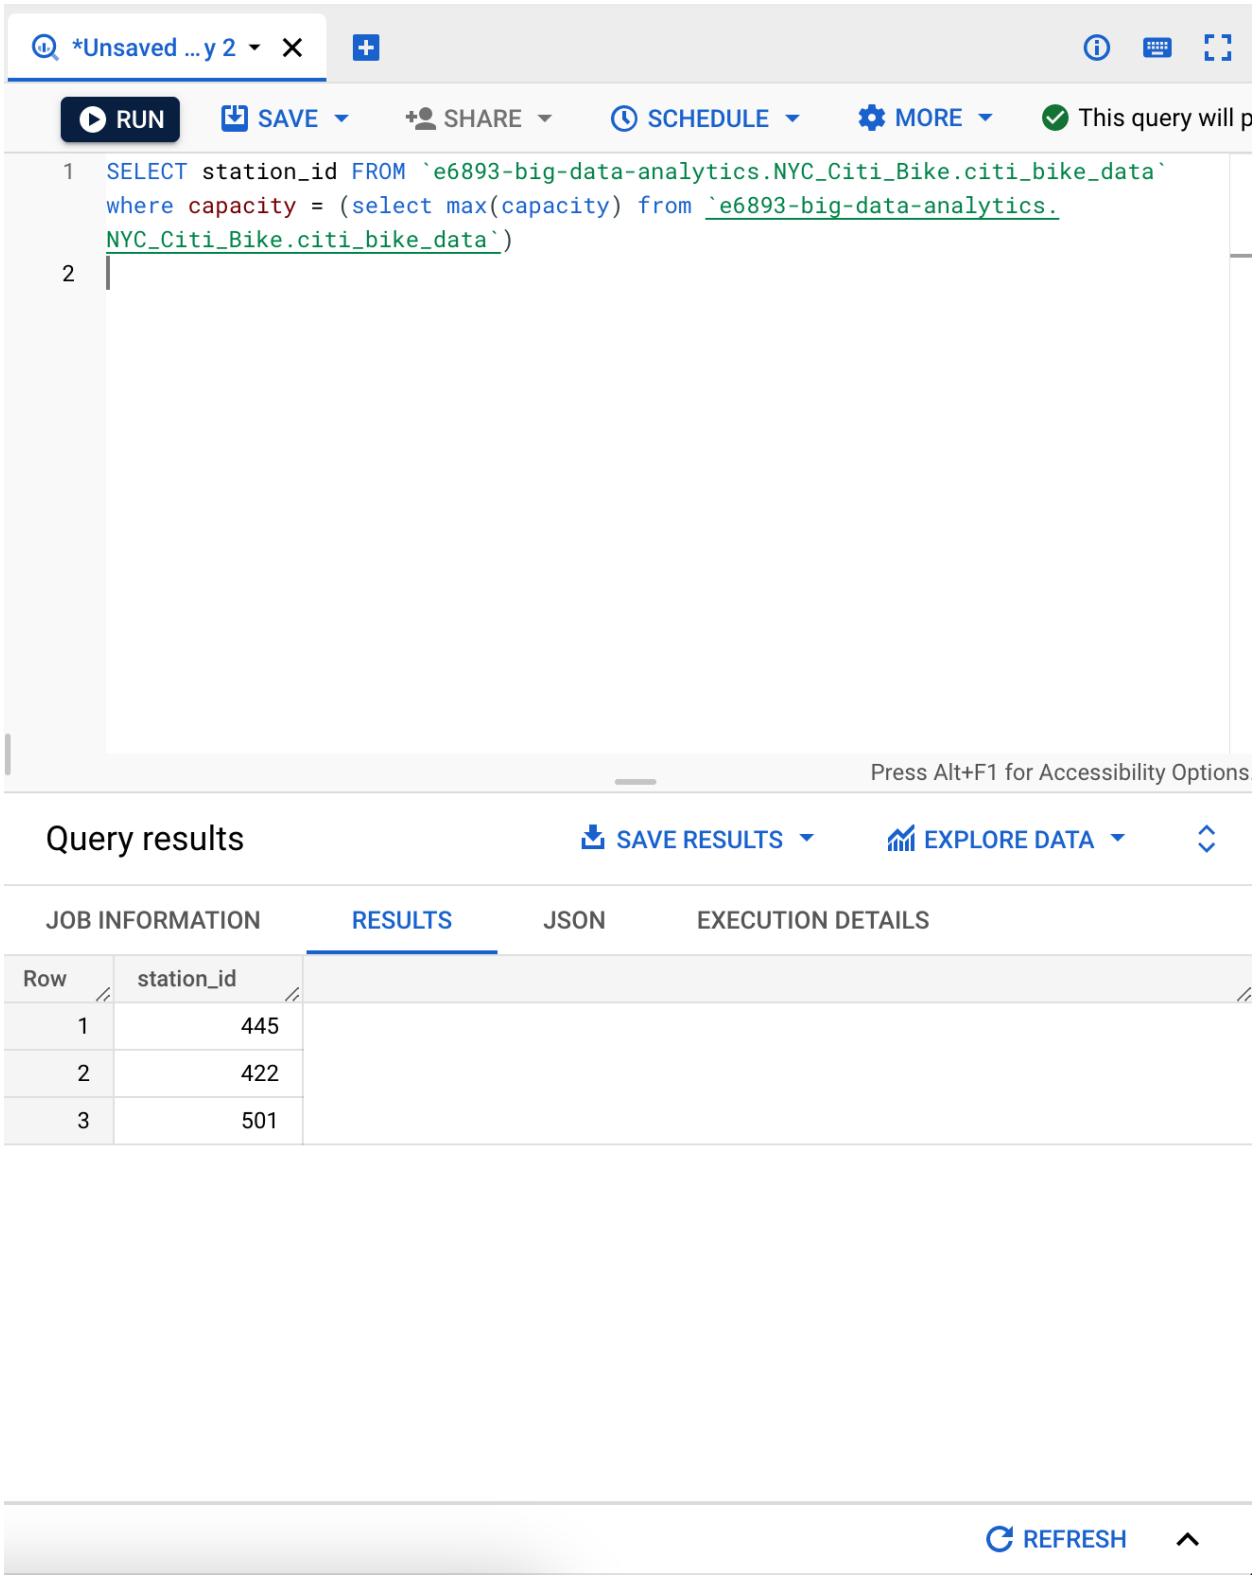
*Unsaved ... y 2
RUN
SAVE
SHARE
SCHEDULE
MORE
This query will p
1 SELECT station_id FROM `e6893-big-data-analytics.NYC_Citi_Bike.citi_bike_data`
where capacity = (select max(capacity) from `e6893-big-data-analytics.
NYC_Citi_Bike.citi_bike_data`)
2
Press Alt+F1 for Accessibility Options.
Query results
SAVE RESULTS
EXPLORE DATA
JOB INFORMATION
RESULTS
JSON
EXECUTION DETAILS
Row
station_id
1 445
2 422
3 501
REFRESH

# Question 3

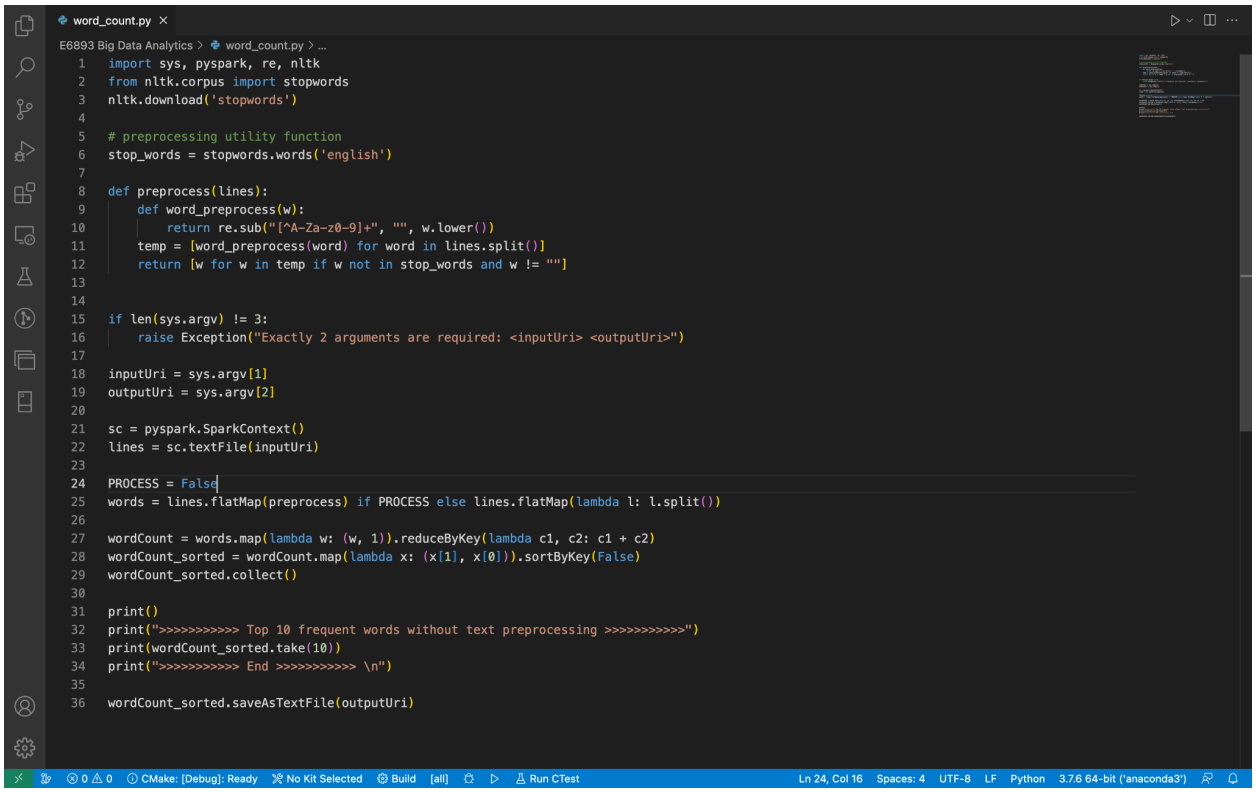


The code above shows the overall workflow of the program. In this exercise, the transformation operations are **‘flatMap’, ‘map’**. The action operations are **‘reduceByKey’, ‘sortByKey’, ‘collect’, ‘take’**.

(1) As the result shows, the top 10 frequent words are ‘the’, ‘and’, ‘of’, ‘to’, ‘I’, ‘a’, ‘you’, ‘in’, ‘is’, ‘my’.

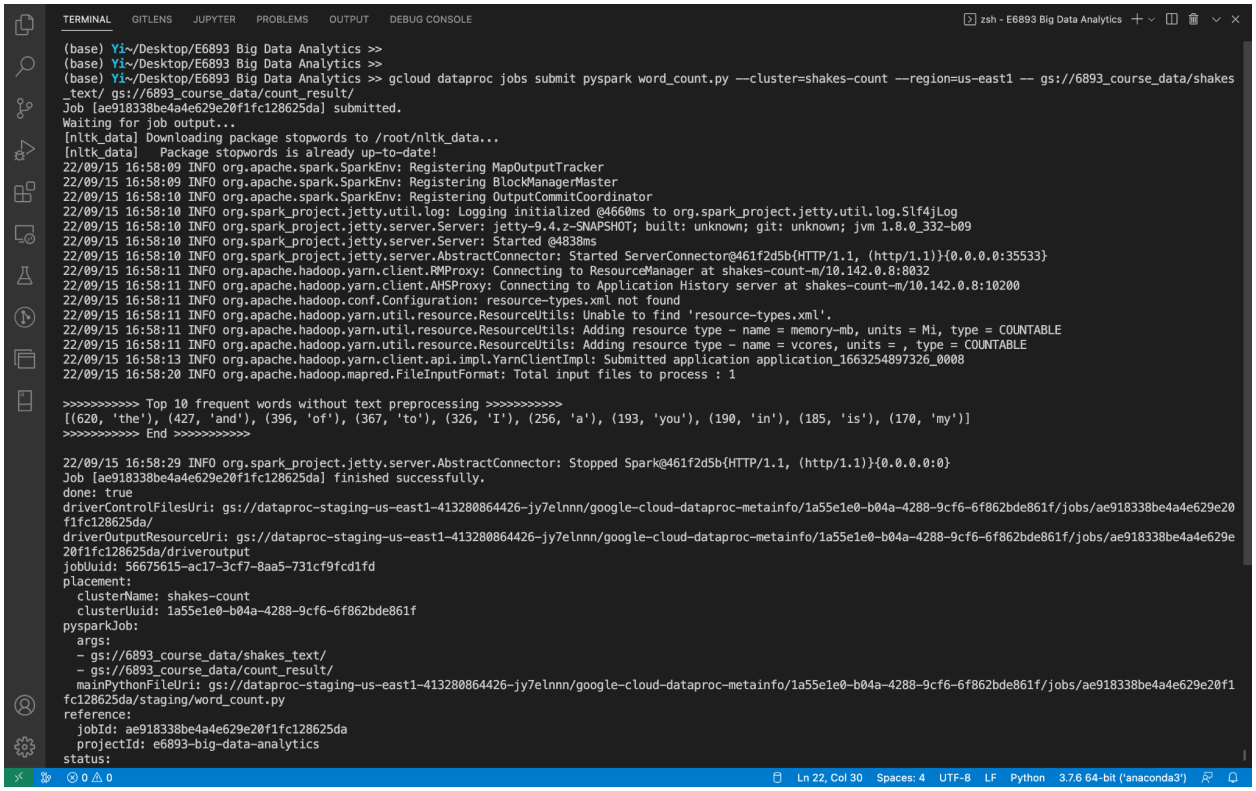


(2) Setting the **PROCESS** variable to **True**, the program will first do preprocessing such as converting all capital letters to lowercase and filtering out stopwords in English. The top 10 after preprocessing becomes **‘macb’, ‘haue’, ‘thou’, ‘enter’, ‘shall’, ‘macbeth’, ‘thee’, ‘vpon’, ‘macd’, ‘yet’**.

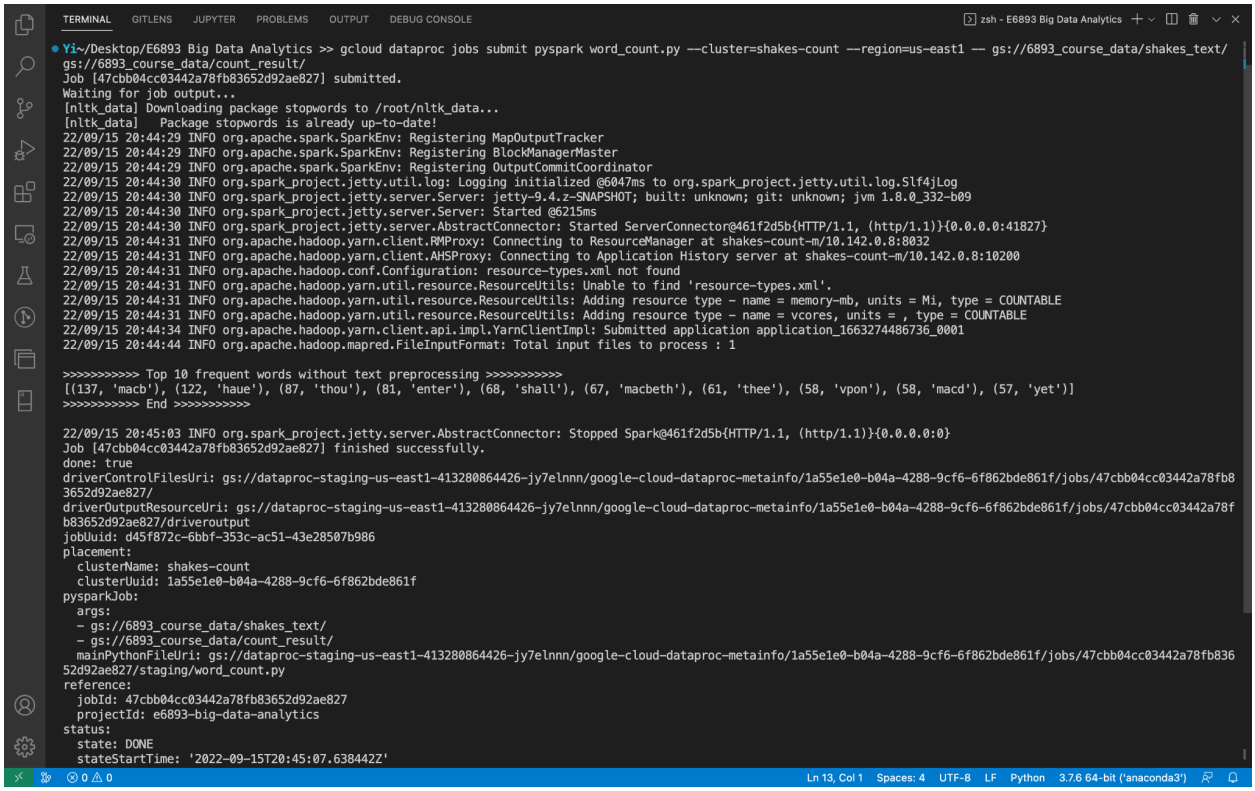
TERMINAL GITLENS JUPYTER PROBLEMS OUTPUT DEBUG CONSOLE
zsh - E6893 Big Data Analytics
Yi~/Desktop/E6893 Big Data Analytics >> gcloud dataproc jobs submit pyspark word_count.py --cluster=shakes-count --region=us-east1 -- gs://6893_course_data/shakes_text/ gs://6893_course_data/count_result/
Job [47cbb04cc03442a78fb83652d92ae827] submitted.
Waiting for job output...
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
22/09/15 20:44:29 INFO org.apache.spark.SparkEnv: Registering MapOutputTracker
22/09/15 20:44:29 INFO org.apache.spark.SparkEnv: Registering BlockManagerMaster
22/09/15 20:44:29 INFO org.apache.spark.SparkEnv: Registering OutputCommitCoordinator
22/09/15 20:44:30 INFO org.spark_project.jetty.util.log: Logging initialized @6047ms to org.spark_project.jetty.util.log.Slf4jLog
22/09/15 20:44:30 INFO org.spark_project.jetty.server.Server: jetty-9.4.z-SNAPSHOT; built: unknown; git: unknown; jvm 1.8.0_332-b09
22/09/15 20:44:30 INFO org.spark_project.jetty.server.Server: Started @6215ms
22/09/15 20:44:30 INFO org.spark_project.jetty.server.AbstractConnector: Started ServerConnector@461f2d5b{HTTP/1.1, (http/1.1)}{0.0.0.0:41827}
22/09/15 20:44:31 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at shakes-count-m/10.142.0.8:8032
22/09/15 20:44:31 INFO org.apache.hadoop.yarn.client.AHSProxy: Connecting to Application History server at shakes-count-m/10.142.0.8:10200
22/09/15 20:44:31 INFO org.apache.hadoop.conf.Configuration: resource-types.xml not found
22/09/15 20:44:31 INFO org.apache.hadoop.yarn.util.resource.ResourceUtils: Unable to find 'resource-types.xml'.
22/09/15 20:44:31 INFO org.apache.hadoop.yarn.util.resource.ResourceUtils: Adding resource type - name = memory-mb, units = Mi, type = COUNTABLE
22/09/15 20:44:31 INFO org.apache.hadoop.yarn.util.resource.ResourceUtils: Adding resource type - name = vcores, units = , type = COUNTABLE
22/09/15 20:44:34 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl: Submitted application application_1663274486736_0001
22/09/15 20:44:44 INFO org.apache.hadoop.mapred.FileInputFormat: Total input files to process : 1
>>>>>>>>>>> Top 10 frequent words without text preprocessing >>>>>>>>>>>
[(137, 'macb'), (122, 'haue'), (87, 'thou'), (81, 'enter'), (68, 'shall'), (67, 'macbeth'), (61, 'thee'), (58, 'vpon'), (58, 'macd'), (57, 'yet')]
>>>>>>>>>>> End >>>>>>>>>>>
22/09/15 20:45:03 INFO org.spark_project.jetty.server.AbstractConnector: Stopped Spark@461f2d5b{HTTP/1.1, (http/1.1)}{0.0.0.0:0}
Job [47cbb04cc03442a78fb83652d92ae827] finished successfully.
done: true
driverControlFilesUri: gs://dataproc-staging-us-east1-413280864426-jy7elnnn/google-cloud-dataproc-metainfo/1a55e1e0-b04a-4288-9cf6-6f862bde861f/jobs/47cbb04cc03442a78fb83652d92ae827/
driverOutputResourceUri: gs://dataproc-staging-us-east1-413280864426-jy7elnnn/google-cloud-dataproc-metainfo/1a55e1e0-b04a-4288-9cf6-6f862bde861f/jobs/47cbb04cc03442a78fb83652d92ae827/driveroutput
jobUuid: d45f872c-6bbf-353c-ac51-43e28507b986
placement:
clusterName: shakes-count
clusterUuid: 1a55e1e0-b04a-4288-9cf6-6f862bde861f
pysparkJob:
args:
- gs://6893_course_data/shakes_text/
- gs://6893_course_data/count_result/
mainPythonFileUri: gs://dataproc-staging-us-east1-413280864426-jy7elnnn/google-cloud-dataproc-metainfo/1a55e1e0-b04a-4288-9cf6-6f862bde861f/jobs/47cbb04cc03442a78fb83652d92ae827/staging/word_count.py
reference:
jobId: 47cbb04cc03442a78fb83652d92ae827
projectId: e6893-big-data-analytics
status:
state: DONE
stateStartTime: '2022-09-15T20:45:07.638442Z'
Ln 13, Col 1 Spaces: 4 UTF-8 LF Python 3.7.6 64-bit ('anaconda3')