



EECS E6893 Big Data Analytics

HW1: Clustering, Spark MLlib, and Hadoop

Tejasri Kurapati, tk2928@columbia.edu

Agenda

- HW1
 - Iterative K-means clustering
 - Hadoop
- Spark MLlib

HW1

HW1

- Document clustering with K-means
 - “Implement” iterative K-means clustering in Spark
 - L1, L2 distance functions
 - Different initialization strategies
 - Plot the cluster assignment result with T-SNE dimensionality reduction
- Monitoring Hadoop metrics
 - Installing Hadoop in Pseudo Distributed Mode
 - Monitoring hadoop metrics through HTTP API

Iterative K-means

- In each iteration, k centroids are initialized, each point in the space is assigned to the nearest centroid, and the centroids are re-computed
- Pseudo code:

Algorithm 1 Iterative k -Means Algorithm

```
1: procedure ITERATIVE  $k$ -MEANS
2:   Select  $k$  points as initial centroids of the  $k$  clusters.
3:   for iterations := 1 to MAX_ITER do
4:     for each point  $p$  in the dataset do
5:       Assign point  $p$  to the cluster with the closest centroid
6:     end for
7:     for each cluster  $c$  do
8:       Recompute the centroid of  $c$  as the mean of all the data points assigned to  $c$ 
9:     end for
10:  end for
11: end procedure
```

Iterative K-means in Spark

Hint:

Spark operations you might need:
map, reduceByKey, collect, keys

Algorithm 1 Iterative k -Means Algorithm

```
1: procedure ITERATIVE  $k$ -MEANS
2:   Select  $k$  points as initial centroids of the  $k$  clusters.
3:   for iterations := 1 to MAX_ITER do
4:     for each point  $p$  in the dataset do
5:       Assign point  $p$  to the cluster with the closest centroid
6:     end for
7:     for each cluster  $c$  do
8:       Recompute the centroid of  $c$  as the mean of all the data points assigned to  $c$ 
9:     end for
10:  end for
11: end procedure
```

```
# iterative k-means
for _ in range(MAX_ITER):
    # Transform each point to a combo of point, closest centroid, count
    # point -> (closest_centroid, (point, 1))

    # Re-compute cluster center

    # For each cluster center (key), aggregate its value by summing up points and count
    # Average the points for each centroid: divide sum of points by count
```

Document clustering

More From Medium

More from Wisecrack



Colonel Sanders Wants to Be Your Daddy



Amanda Sche... Wisecrack
Sep 13 · 5 min read ★



132



More from Wisecrack



Platforms Own You, Now What?



Wisecrack in Wisecrack
Sep 10 · 6 min read ★



590



More from Wisecrack



Millennials Are Shaking Up Wall Street, But Is It Helping Them?



Thomas Ambr... Wisecrack
Sep 12 · 9 min read ★



14



Two reasons why Antoine Griezmann has doubts over moving to Man United

independent.co.uk · 1h · Miguel Delaney an...



Azuz11

added to Soccer Player

Write a comment...



MORE STORIES:

Manchester United:
Jose Mourinho £85
Million Transfer
Target 'Agrees Per...

NEWSWEEK.COM

Mar
secu
with
Grie

INTE



COVER STORIES

Plot the result with t-SNE

```
from sklearn.manifold import TSNE
```

```
# RDD -> np array  
data_np = np.array(data.collect())
```

```
data_np.shape
```

```
(4601, 58)
```


```
data_embedded = TSNE(n_components=2).fit_transform(data_np)
```

```
data_embedded.shape
```


```
(4601, 2)
```

```
vis_x = data_embedded[:, 0]  
vis_y = data_embedded[:, 1]  
plt.scatter(vis_x, vis_y, cmap=plt.cm.get_cmap("jet", 10))  
plt.show()
```


Plot the result with t-SNE




Before clustering




After clustering

Plot the result with t-SNE (set random state)



t-SNE results of c1.txt using L2 distance and random_state=100



t-SNE results of c2.txt using L2 distance and random_state=100



Plot the cost of each iteration



Spark MLlib

- Spark's scalable machine learning library
- Tools:
 - ML Algorithms: classification, regression, clustering, and collaborative filtering
 - Featurization: feature extraction, transformation, dimensionality reduction, and selection
 - Pipelines: tools for constructing, evaluating, and tuning ML Pipelines
 - Persistence: saving and load algorithms, models, and Pipelines
 - Utilities: linear algebra, statistics, data handling, etc.

Example: K-means clustering with Spark MLlib

```
from pyspark.mllib.clustering import KMeans
```

```
clusters = KMeans.train(data, 10, maxIterations=20, initializationMode="random")
```

```
# cluster centers  
len(clusters.centers)
```

10

Hadoop installation

Step 1: Pre-installation Setup

- Before the installation, learn how to login & exit the root account
 - Login: `sudo -i`
 - Exit: `exit` (or use `ctrl+D`)

```
(base) yl@Yvonne-surfacebook2:/mnt/c/Users/sh_yv$ sudo -i
[sudo] password for yl:
root@Yvonne-surfacebook2:~# exit
logout
(base) yl@Yvonne-surfacebook2:/mnt/c/Users/sh_yv$ |
```

● Create a user

- Open the **root** using the command “`sudo -i`”.
- Create a user from the root account using the command “`useradd -m username`”.
- Set the password using the command “`passwd username`”
- Now you can open the new user account.
 - If you're under root account, use the command “`su username`”
 - Otherwise, use “`su - username`”

```
(base) yl@Yvonne-surfacebook2:/mnt/c/Users/sh_yv$ sudo -i
[sudo] password for yl:
root@Yvonne-surfacebook2:~# XXXXXXXXXX
root@Yvonne-surfacebook2:~# useradd -m hadoop
root@Yvonne-surfacebook2:~# passwd hadoop
New password:
Retype new password:
passwd: password updated successfully
```


- Create a user
 - Add user to sudo group

```
root@Yvonne-surfacebook2:~# adduser hadoop sudo
Adding user `hadoop' to group `sudo' ...
Adding user hadoop to group sudo
Done.
```

● SSH Setup and Key Generation

- Open the account you created, using
 - `su hadoop`
- Generate generating a key value pair using SSH, using
 - `ssh-keygen -t rsa` (press “enter” directly where you’re asked to enter)
- Copy the public keys from `id_rsa.pub` to `authorized_keys`, using
 - `cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys`
- Provide the owner with read and write permissions to `authorized_keys` file respectively
 - `chmod 0600 ~/.ssh/authorized_keys`
- Test SSH setup
 - `ssh localhost`

```
root@Yvonne-surfacebook2:~# su hadoop
$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/hadoop/.ssh/id_rsa):
Created directory '/home/hadoop/.ssh'.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/hadoop/.ssh/id_rsa
Your public key has been saved in /home/hadoop/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:gl3Zvwde00N6gVncTWlUyc22YoKFP4HsuhfhVmIMRJY hadoop@Yvonne-surfacebook2
The key's randomart image is:
+---[RSA 3072]-----+
|      o+o o   oo*|
|      .E * o   ==|
|      * =....+.|
|      o . B *o+...|
|      . o S ++X o|
|      o +oo.*|
|      o . +..|
|      . . . .|
|      . . . .|
+-----[SHA256]-----+
$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
$ chmod 0600 ~/.ssh/authorized_keys
```

- Test ssh setup. Use “logout” command to log out

```
$ ssh localhost
Welcome to Ubuntu 20.04.2 LTS (GNU/Linux 5.4.72-microsoft-standard-WSL2 x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

System information as of Thu Sep 23 15:37:35 EDT 2021

System load:  0.0               Processes:            24
Usage of /:   11.0% of 250.98GB Users logged in:      0
Memory usage: 1%               IPv4 address for eth0: 172.19.193.5
Swap usage:   0%

213 updates can be installed immediately.
91 of these updates are security updates.
To see these additional updates run: apt list --upgradable

Last login: Thu Sep 23 15:32:08 2021 from 127.0.0.1
```

- SSH Setup (for Debugging)

- If ssh localhost doesn't work

```
$ ssh localhost  
ssh: connect to host localhost port 22: Connection refused
```

- Try reinstall some packages:
 - `sudo apt-get remove openssh-client openssh-server`
 - `sudo apt-get install openssh-client openssh-server`
- If still doesn't work, check the following
 - `sudo service ssh start`
 - `ssh localhost`

● Installing Java

- Verify the existence of Java in your system
 - `java -version`
 - If you've installed Java, it will give you the following output, and you can skip the java installing steps, continuing to the next section.

```
java version "1.7.0_71"  
Java(TM) SE Runtime Environment (build 1.7.0_71-b13)  
Java HotSpot(TM) Client VM (build 25.0-b02, mixed mode)
```

- If you did not install Java, you need to follow the next steps to install Java.

● Installing Java

- Install java
 - `sudo apt-get install openjdk-8-jre openjdk-8-jdk`
- Then check Java version to see if you have installed java
 - `java -version`
- To find where you have installed java
 - `dirname $(dirname $(readlink -f $(which javac)))`

```
(base) yl@Yvonne-surfacebook2:/usr/bin$ dirname $(dirname $(readlink -f $(which javac)))  
/usr/lib/jvm/java-8-openjdk-amd64
```

- Set up PATH and JAVA_HOME variables
 - `export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64` (the path from last step)
 - `export PATH=$PATH:$JAVA_HOME/bin`
- Now apply all the changes into the current running system.
 - `exec bash`

Step 2: Downloading Hadoop

- Change to root and change directory
 - `sudo -i`
 - `cd /usr/local/`
- Download and extract Hadoop
 - `wget http://apache.claz.org/hadoop/common/hadoop-3.3.1/hadoop-3.3.1.tar.gz`
 - `tar xzf hadoop-3.3.1.tar.gz`
 - `mv hadoop-3.3.1 hadoop`
- Change owner
 - `sudo chown -R hadoop:hadoop ./hadoop`
- Set Hadoop environment variables
 - `su hadoop`
 - `export HADOOP_HOME=/usr/local/hadoop`
 - `export PATH=$PATH:/usr/local/hadoop/bin`
 - `exec bash`

- Test the Hadoop setup

- Type the following command
 - `hadoop version`
 - If everything is fine, you'll see the following

```
$ hadoop version
Hadoop 3.3.1
Source code repository https://github.com/apache/hadoop.git -r a3b9c37a397ad4188041dd80621bdeefc46885f2
Compiled by ubuntu on 2021-06-15T05:13Z
Compiled with protoc 3.7.1
From source with checksum 88a4ddb2299aca054416d6b7f81ca55
This command was run using /usr/local/hadoop/share/hadoop/common/hadoop-common-3.3.1.jar
```

- Now you have successfully set up the Hadoop's standalone mode

● Installing Hadoop in Pseudo Distributed Mode

- Set the Hadoop environment variables
 - `export HADOOP_HOME=/usr/local/hadoop`
 - `export HADOOP_MAPRED_HOME=$HADOOP_HOME`
 - `export HADOOP_COMMON_HOME=$HADOOP_HOME`
 -
 - `export HADOOP_HDFS_HOME=$HADOOP_HOME`
 - `export YARN_HOME=$HADOOP_HOME`
 - `export`
`HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native`
 - `export PATH=$PATH:$HADOOP_HOME/sbin:$HADOOP_HOME/bin`
 - `export HADOOP_INSTALL=$HADOOP_HOME`
 - `exec bash`

- Hadoop configuration

- Find the Hadoop configuration files

- `cd $HADOOP_HOME/etc/hadoop`
- `vim hadoop-env.sh` (Add the location of java to this file, namely the following line)
- `JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64`

```
###  
  
#  
# When building Hadoop, one can add the class paths to the commands  
# via this special env var:  
# export HADOOP_ENABLE_BUILD_PATHS="true"  
  
#  
# To prevent accidents, shell commands be (superficially) locked  
# to only allow certain users to execute certain subcommands.  
# It uses the format of (command)_(subcommand)_USER.  
#  
# For example, to limit who can execute the namenode command,  
# export HDFS_NAMENODE_USER=hdfs  
  
###  
# Registry DNS specific parameters  
###  
# For privileged registry DNS, user to run as after dropping privileges  
# This will replace the hadoop.id.str Java property in secure mode.  
# export HADOOP_REGISTRYDNS_SECURE_USER=yarn  
  
# Supplemental options for privileged registry DNS  
# By default, Hadoop uses jsvc which needs to know to launch a  
# server jvm.  
# export HADOOP_REGISTRYDNS_SECURE_EXTRA_OPTS="--jvm server"  
JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64  
"hadoop-env.sh" 431L, 16698C
```

- Some files that you need to edit to configure Hadoop

- Open the core-site.xml and add the following properties in between `<configuration>`, `</configuration>` tags.

```
<configuration>
```

```
  <property>
```

```
    <name>fs.default.name</name>
```

```
    <value>hdfs://localhost:9000</value>
```

```
  </property>
```

```
</configuration>
```

- Some files that you need to edit to configure Hadoop

- Open the hdfs-site.xml and add the following properties in between <configuration>, </configuration> tags.

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>

  <property>
    <name>dfs.name.dir</name>
    <value>file:///home/hadoop/hadoopinfra/hdfs/namenode </value>
  </property>

  <property>
    <name>dfs.data.dir</name>
    <value>file:///home/hadoop/hadoopinfra/hdfs/datanode </value>
  </property>
</configuration>
```

- Some files that you need to edit to configure Hadoop

- Open the yarn-site.xml and add the following properties in between <configuration>, </configuration> tags.

```
<configuration>
```

```
  <property>
```

```
    <name>yarn.nodemanager.aux-services</name>
```

```
    <value>mapreduce_shuffle</value>
```

```
  </property>
```

```
</configuration>
```

- Some files that you need to edit to configure Hadoop

- Open the mapred-site.xml and add the following properties in between <configuration>, </configuration> tags.

```
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>
```


● Verify Hadoop installation

- Set up the namenode
 - `cd ~`
 - `hdfs namenode -format`

```
10/24/14 21:30:55 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG:  host = localhost/192.168.1.11
STARTUP_MSG:  args = [-format]
STARTUP_MSG:  version = 2.4.1
...
...
10/24/14 21:30:56 INFO common.Storage: Storage directory
/home/hadoop/hadoopinfra/hdfs/namenode has been successfully formatted.
10/24/14 21:30:56 INFO namenode.NNStorageRetentionManager: Going to
retain 1 images with txid >= 0
10/24/14 21:30:56 INFO util.ExitUtil: Exiting with status 0
10/24/14 21:30:56 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at localhost/192.168.1.11
*****/
```

- Verify Hadoop installation

- Verify Hadoop dfs

- Start-dfs.sh

```
hadoop@Yvonne-surfacebook2:/usr/local/hadoop/etc/hadoop$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [Yvonne-surfacebook2]
```

- Verify yarn script

- start-yarn.sh

```
hadoop@Yvonne-surfacebook2:~$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
```

● Access Hadoop on Browser

- Use the following url to get Hadoop services on browser.
 - <http://localhost:9870/>



Overview 'localhost:9000' (✓active)

Started:	Thu Sep 23 17:49:18 -0400 2021
Version:	3.3.1, ra3b9c37a397ad4188041dd80621bdeefc46885f2
Compiled:	Tue Jun 15 01:13:00 -0400 2021 by ubuntu from (HEAD detached at release-3.3.1-RC3)
Cluster ID:	CID-a9dc07bc-f1c1-4847-877c-78f23e32729a
Block Pool ID:	BP-1874079373-127.0.1.1-1632432704363

Summary

Security is off.
Safemode is off.

● Access Hadoop on Browser

- Access all applications of cluster
 - <http://localhost:8088/>



Cluster

[About](#)
[Nodes](#)
[Node Labels](#)
[Applications](#)
[NEW](#)
[NEW SAVING](#)
[SUBMITTED](#)
[ACCEPTED](#)
[RUNNING](#)
[FINISHED](#)
[FAILED](#)
[KILLED](#)
[Scheduler](#)

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	
0	0	0	0	0	<memory>

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes
1	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>

Show 20 entries

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State
Showing 0 to 0 of 0 entries										

References

- <https://spark.apache.org/docs/latest/sql-getting-started.html>
- <https://www.analyticsvidhya.com/blog/2016/10/spark-dataframe-and-operation-s/>
- <https://spark.apache.org/docs/latest/ml-guide.html>
- <https://towardsdatascience.com/machine-learning-with-pyspark-and-mllib-solving-a-binary-classification-problem-96396065d2aa>